

Search trees, binary trie, patricia trie

Marko Berezovský
Radek Mařík
PAL 2012

To read

Robert Sedgwick : *Algorithms in C++, Parts 1–4: Fundamentals, Data Structure, Sorting, Searching*, Third Edition, Addison Wesley 1998, chapter 15.

See PAL webpage for references

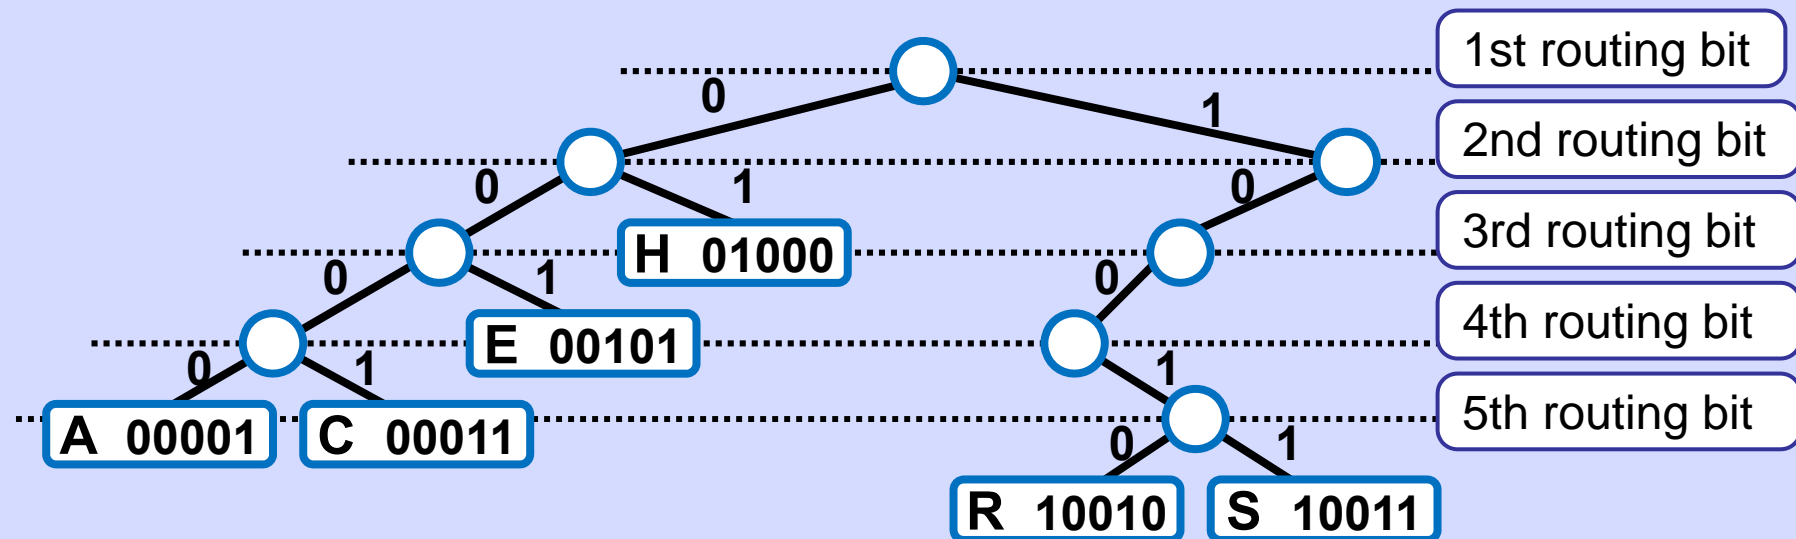
Keys are represented as a sequence of bits.

Keys are stored in roots only, inner nodes serve as routers only.


Inner node in depth d defines path to the leaf with key K according to the d -th bit of K .

Key must not be prefix of another key in terms of bit representation.


This can be achieved by representations having all same length in bits.



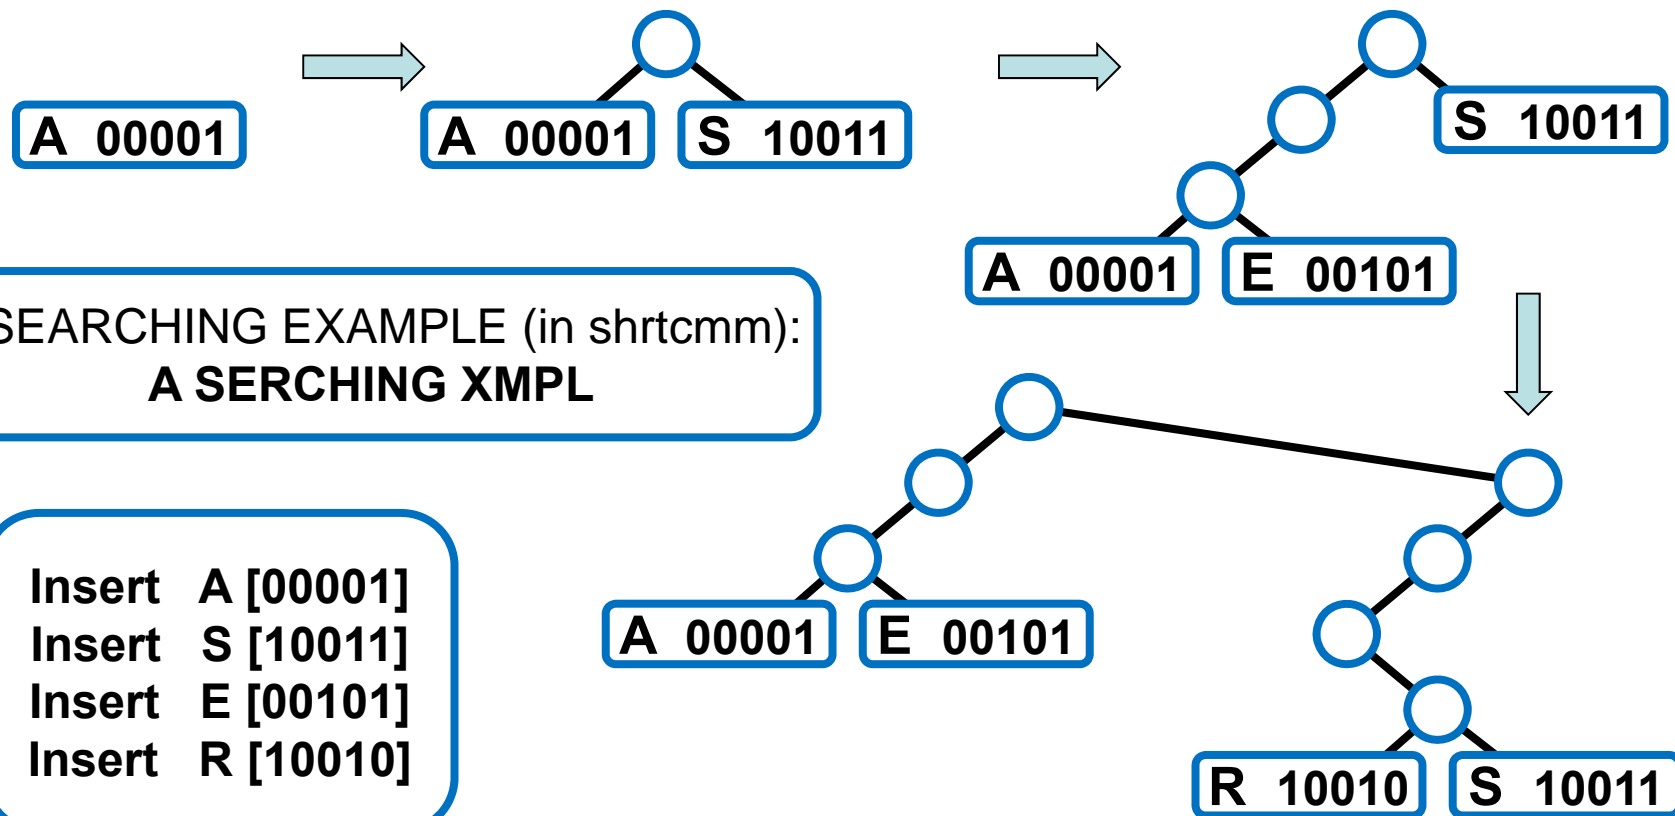
Tree height \leq no of routing bits = length of bit representation of keys



Inserting single key may result in creating more internal nodes. Insert I [01001] .



Example of trie building.




A SEARCHING EXAMPLE (in shrtcmm):
A SERCHING XMPL


Insert A [00001]
Insert S [10011]
Insert E [00101]
Insert R [10010]

In further examples we omit 0/1 edge labels as the labeling scheme is obvious.


Insert C [00011]




Insert H [01000]




Insert I [01001]




Insert N [01110]



Insert G [00111]





Null pointers in internal nodes are called null links. When trie has no null link it is a complete perfectly balanced binary tree containing $2^{d+1}-1$ nodes and 2^d leaves.

The function `searchR` uses the bits of the key to control the branching on the way down the trie.

There are three possible outcomes: if the search reaches a leaf (with both links null), then that is the unique node in the trie that could contain the record with key v , so we test whether that node indeed contains v (search hit) or some key whose leading bits match v (search miss). If the search reaches a null link, then the parent's other link must not be null, so there is some other key in the trie that differs from the search key in the corresponding bit, and we have a search miss.

The following code assumes that the keys are distinct, and (if the keys may be of different lengths) that no key is a prefix of another. The `item` member is not used in non-leaf nodes.

private:

```
Item searchR(link h, Key v, int d)
{ if (h == 0) return nullItem;
  if (h->l == 0 && h->r == 0)
  { Key w = h->item.key();
    return (v == w) ? h->item : nullItem; }
  if (digit(v, d) == 0)
    return searchR(h->l, v, d+1);
  else return searchR(h->r, v, d+1);
}
```

public:

```
Item search(Key v)
{ return searchR(head, v, 0); }
```

To insert a new node into a trie, we search as usual, then distinguish the two cases that can occur for a search miss.

If the miss was not on a leaf, then we replace the null link that caused us to detect the miss with a link to a new node, as usual.

If the miss was on a leaf, then we use a function `split` to make one new internal node for each bit position where the search key and the key found agree, finishing with one internal node for the leftmost bit position where the keys differ.

The switch statement in `split` converts the two bits that it is testing into a number to handle the four possible cases.

If the bits are the same (case $00_2 = 0$ or $11_2 = 3$), then we continue splitting; if the bits are different (case $01_2 = 1$ or $10_2 = 2$), then we stop splitting.


private:

```
link split(link p, link q, int d)
{ link t = new node(nullItem); t->N = 2;
  Key v = p->item.key(); Key w = q->item.key();
  switch (digit(v, d)*2 + digit(w, d))
  { case 0: t->l = split(p, q, d+1); break;
    case 1: t->l = p; t->r = q; break;
    case 2: t->r = p; t->l = q; break;
    case 3: t->r = split(p, q, d+1); break;
  }
  return t;
}

void insertR(link& h, Item x, int d)
{ if (h == 0) { h = new node(x); return; }
  if (h->l == 0 && h->r == 0)
  { h = split (new node(x), h, d); return; }
  if (digit(x.key(), d) == 0)
    insertR(h->l, x, d+1);
  else insertR(h->r, x, d+1);
}
```

```
public:
    ST(int maxN)
        { head = 0; }
    void insert(Item item)
        { insertR(head, item, 0); }
```

This trie, schematically presented here, built by inserting about 200 random keys, is well-balanced, but has 44 percent more nodes than might otherwise be necessary, because of one-way branching.



Operation Delete is complementary to Insert operation. Its algorithm/implementation is left to the reader.

The structure of a trie is independent of the key insertion order: There is a unique trie for any given set of distinct keys.

Insertion or search for a random key in a trie built from N random (distinct) bitstrings requires about $\lg N$ bit comparisons on the average. The worst-case number of bit comparisons is bounded only by the number of bits in the search key.

A trie built from N random w -bit keys has about $N / \ln 2 \approx 1.44N$ nodes on the average.

P	A	T	R	I	C	I	A
r	l	o	e	n	o	n	l
a	g		t	f	d		p
c	o		r	o	e		h
t	r		i	r	d		a
i	i		e	m			n
c	t		v	a			u
a	h		e	t			m
l	m			i			e
				o			r
				n			i
							c

Donal R. Morrison: *PATRICIA—Practical Algorithm To Retrieve Information Coded in Alphanumeric*, Journal of the ACM, Volume 15 Issue 4, Oct. 1968, pp 514-534.

Trie built by inserting about 200 random keys



Patricia trie built by inserting about 200 random keys



Compare how well are both trees balanced

Starting with the standard trie data structure, we avoid one-way branching via a simple device: we put into each node the index of the bit to be tested to decide which path to take out of that node.

Thus, we jump directly to the bit where a significant decision is to be made, bypassing the bit comparisons at nodes where all the keys in the subtree have the same bit value.

Moreover, we avoid external nodes via another simple device: we store data in internal nodes and replace links to external nodes with links that point back upwards to the correct internal node in the trie.

These two changes allow us to represent tries with binary trees comprising nodes with a key and two links (and an additional field for the index), which we call patricia tries. With patricia tries, we store keys in nodes as with tries, and we traverse the tree according to the bits of the search key, but we do not use the keys in the nodes on the way down the tree to control the search; we merely store them there for possible later reference, when the bottom of the tree is reached.

This recursive procedure shows the records in a patricia trie in order of their keys. We imagine the items to be in (virtual) external nodes, which we can identify by testing when the bit index on the current node is not larger than the bit index on its parent. Otherwise, this program is a standard inorder traversal.

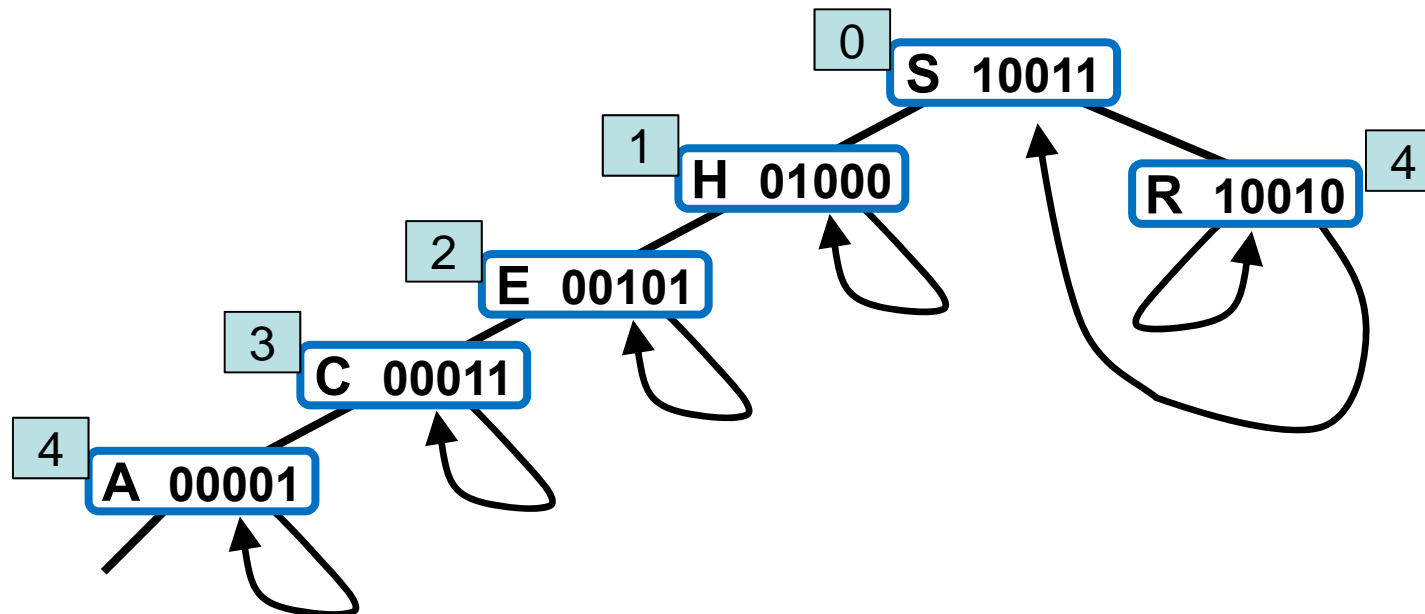
private:

```
void showR(link h, ostream& os, int d)
{
    if (h->bit <= d) { h->item.show(os); return; }
    showR(h->l, os, h->bit);
    showR(h->r, os, h->bit);
}
```

public:

```
void show(ostream& os)
{ showR(head->l, os, -1); }
```

In a successful search for $R = 10010$ in this sample patricia trie (top), we move right (since bit 0 is 1), then left (since bit 4 is 0), which brings us to R (the only key in the tree that begins with $1^{***}0$). On the way down the tree, we check only the key bits indicated in the numbers over the nodes (and ignore the keys in the nodes). When we first reach a link that points up the tree, we compare the search key against the key in the node pointed to by the up link, since that is the only key in the tree that could be equal to the search key.



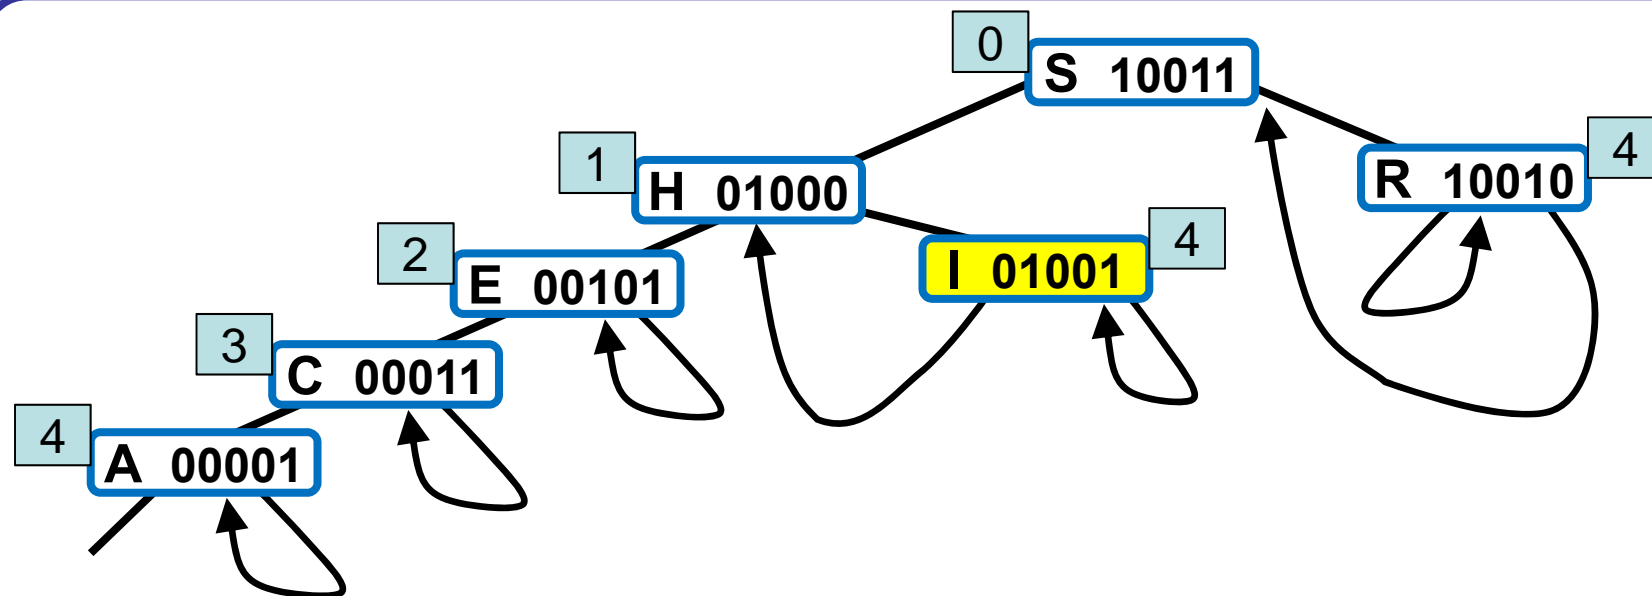
In an unsuccessful search for $I = 01001$, we move left at the root (since bit 0 of the key is 0), then take the right (up) link (since bit 1 is 1) and find that H (the only key in the trie that begins with 01) is not equal to I .

The Search method differs from trie search in three ways: there are no explicit null links, we test the indicated bit in the key instead of the next bit, and we end with a search key comparison at the point where we follow a link up the tree. It is easy to test whether a link points up, because the bit indices in the nodes (by definition) increase as we travel down the tree. To search, we start at the root and proceed down the tree, using the bit index in each node to tell us which bit to examine in the search key—we go right if that bit is 1, left if it is 0. The keys in the nodes are not examined at all on the way down the tree. Eventually, an upward link is encountered: each upward link points to the unique key in the tree that has the bits that would cause a search to take that link. Thus, if the key at the node pointed to by the first upward link encountered is equal to the search key, then the search is successful; otherwise, it is unsuccessful.

The recursive function `searchR` returns the unique node that could contain the record with key `v`. It travels down the trie, using the bits of the tree to control the search, but tests only 1 bit per node encountered—the one indicated in the bit field. It terminates the search when it encounters an external link, one which points up the tree. The search function `search` calls `searchR`, then tests the key in that node to determine whether the search is a hit or a miss.

```
private:
    Item searchR(link h, Key v, int d) {
        if (h->bit <= d) return h->item;
        if (digit(v, h->bit) == 0)
            return searchR(h->l, v, h->bit);
        else return searchR(h->r, v, h->bit);
    }
public:
    Item search(Key v)
    { Item t = searchR(head, v, -1);
      return (v == t.key()) ? t : nullItem;
    }
```

To insert I into the sample patricia trie in , we add a new node to check bit 4, since $H = 01000$ and $I = 01001$ differ in only that bit (top). On a subsequent search in the trie that comes to the new node, we want to check H (left link) if bit 4 of the search key is 0; if the bit is 1 (right link), the key to check is I.



To insert $N = 01110$, we add a new node in between H and I to check bit 2, since that bit distinguishes N from H and I .

